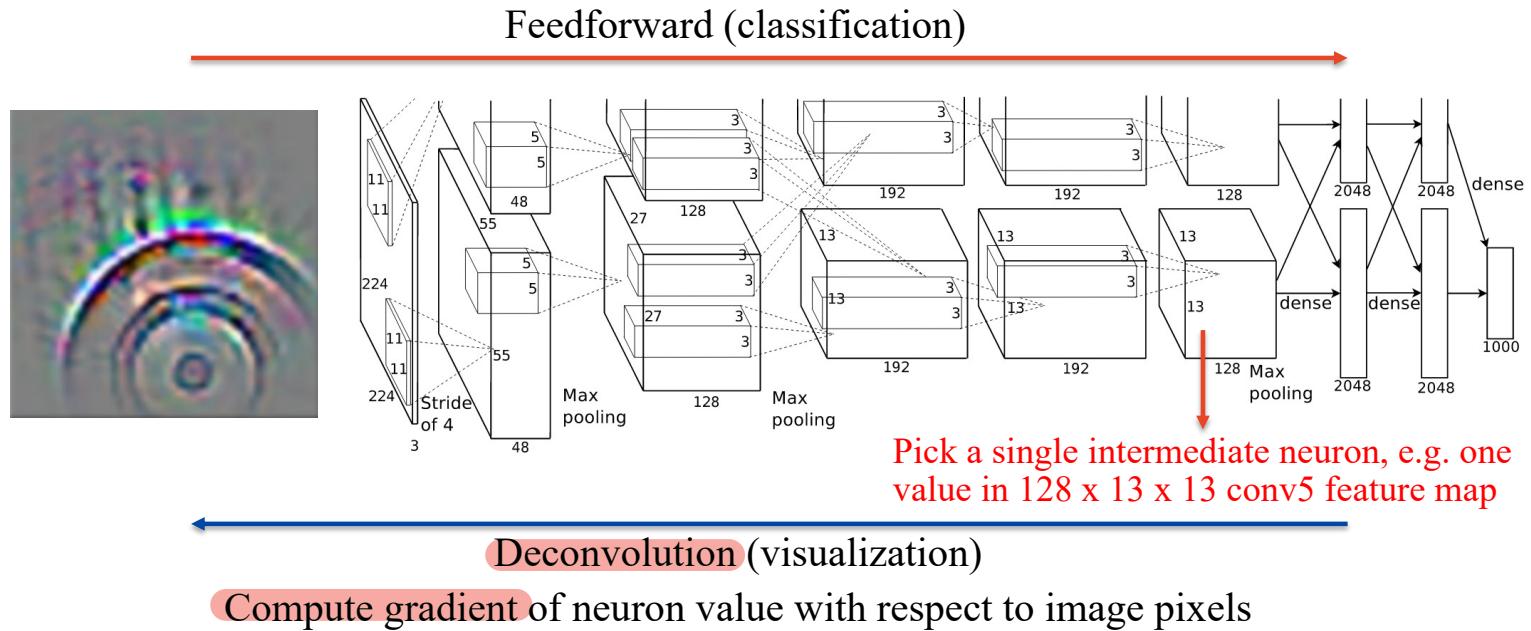


Extension

Visualize features

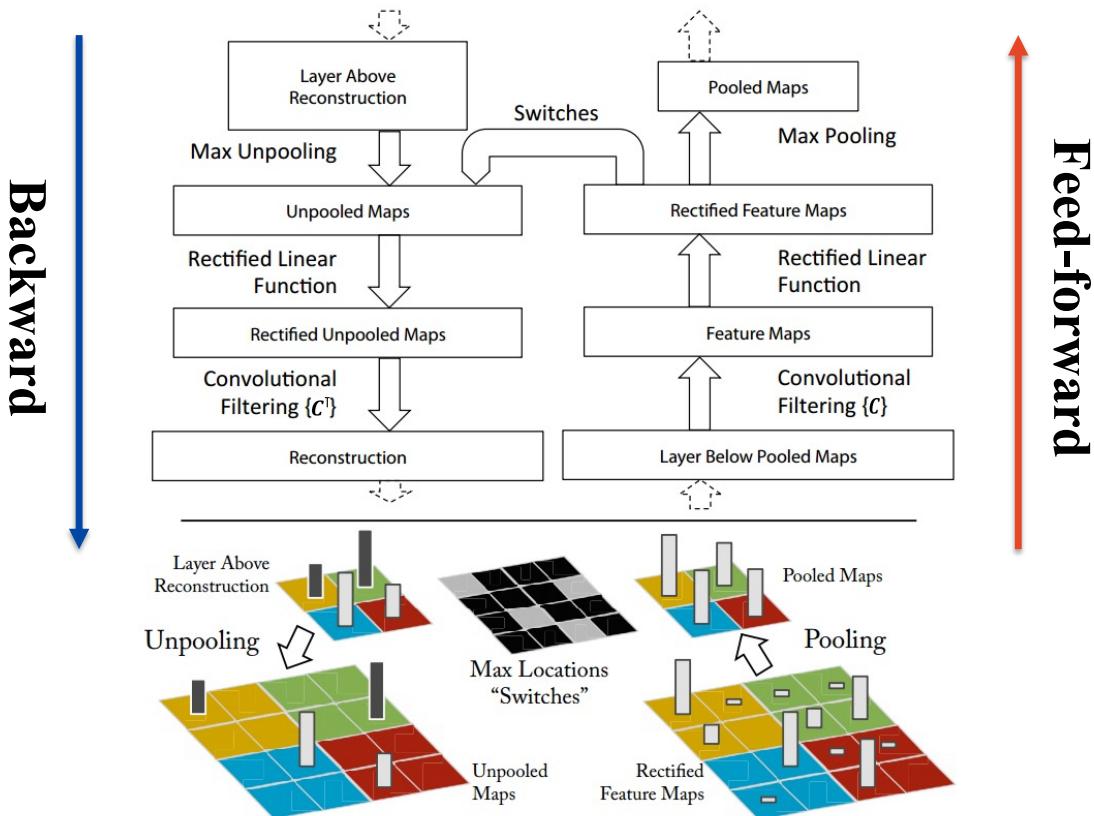
Zeiler and Fergus, ECCV 2014

- A non-parametric view of invariance, showing which patterns from the training set activate the feature map
- Using deconvolution to project the feature activations back to the input pixel space



Deconvolutional Method

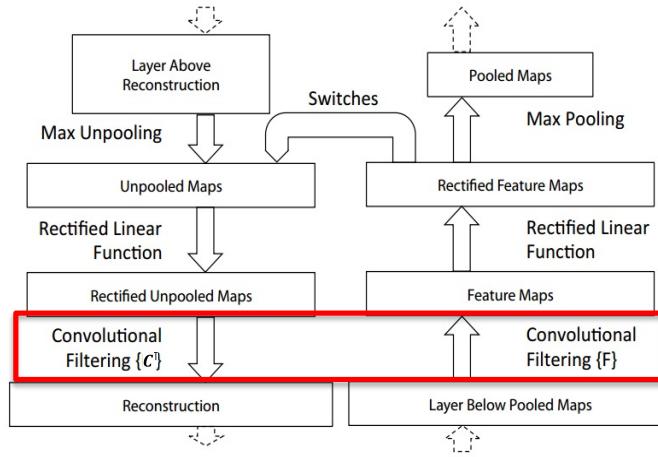
- For one convolutional layer.



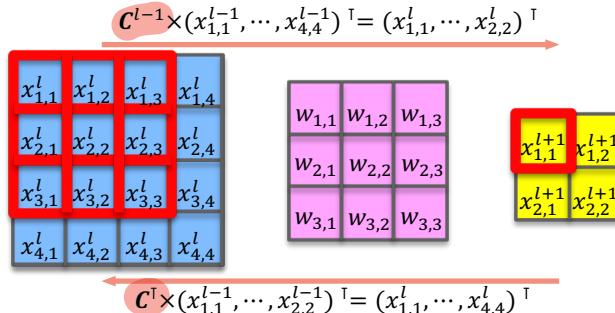
(Zeiler and Fergus, 2014)

Deconvolutional Method

- Conv and Deconv
 - Use learned filters $\{C\}$ to convolve forward signal, but use transposed versions $\{C^\top\}$ for backward signal.

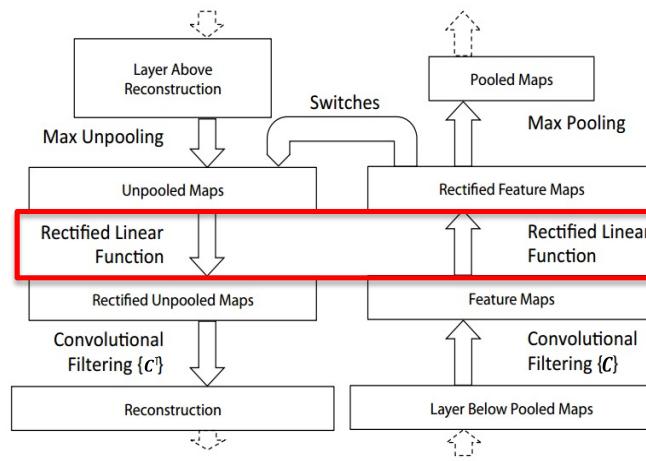


(Zeiler and Fergus, 2014)



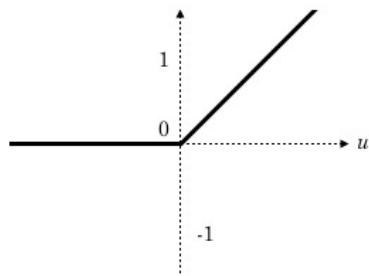
Deconvolutional Method

- Activation function
 - Both feedforward and backward processes employ ReLU as activation functions to ensure feature maps are **always positive**.



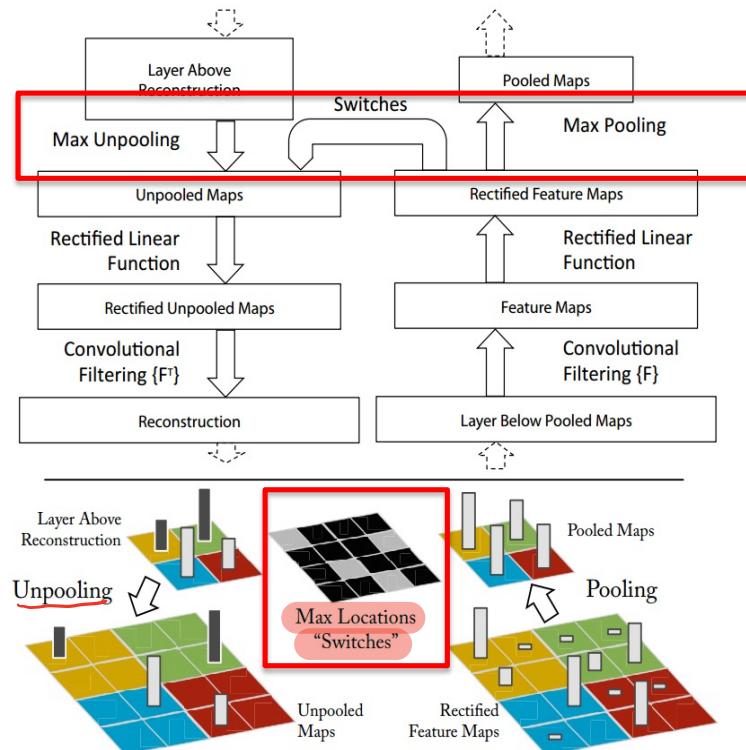
(Zeiler and Fergus, 2014)

$$ReLU(x) = \max(0, x)$$



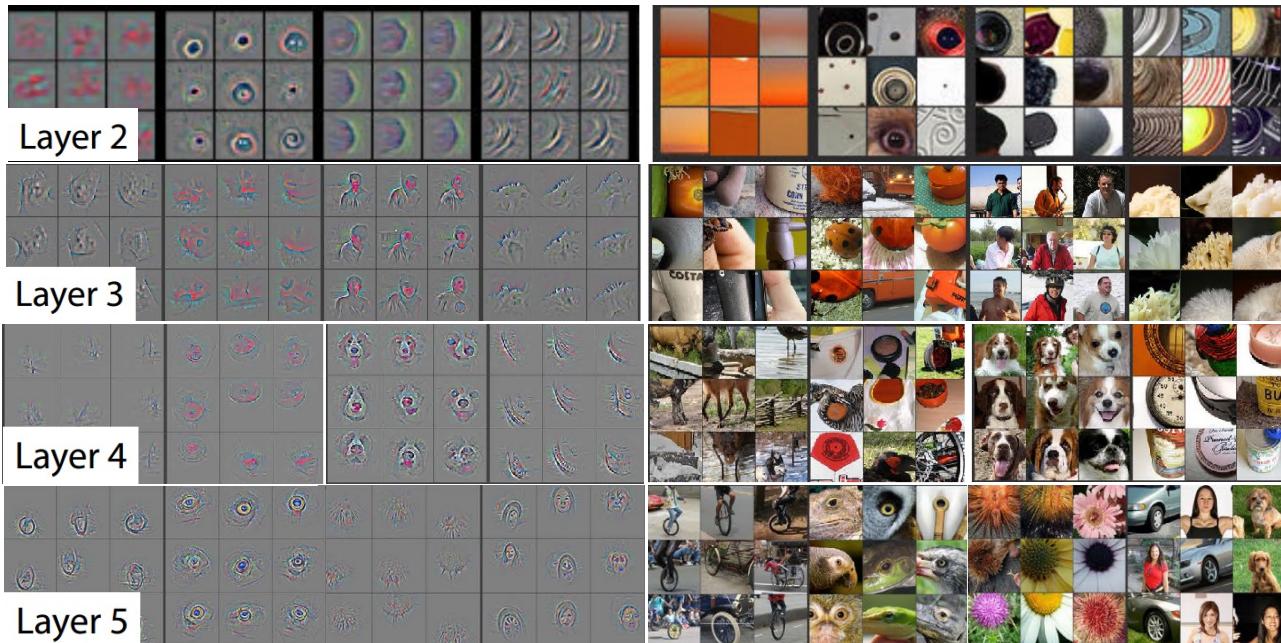
Deconvolutional Method

- Pooling and Unpooling
 - Despite the non-invertibility of max pooling, an approximate inverse could be obtained by recording the locations of the maxima within each pooling region in a set of **switch variables**.



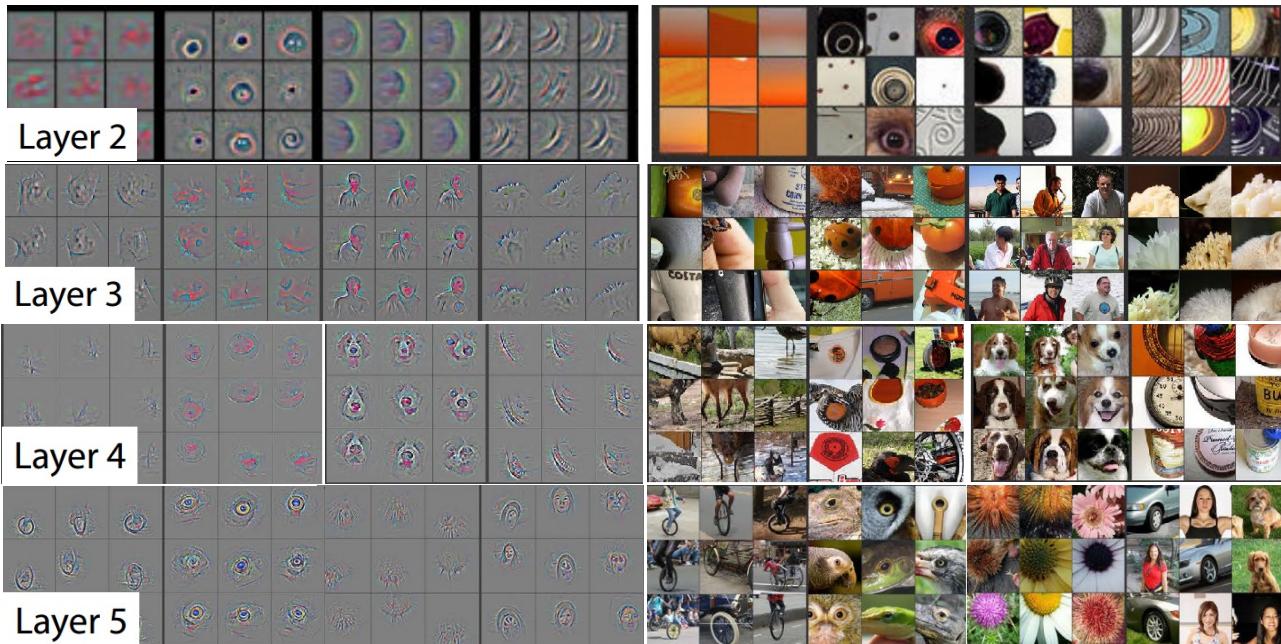
Visualize features

- Instead of showing the single strongest activation for a given feature map, the article shows the top 9 activations layer 2-5.
- Alongside these visualizations are the corresponding image patches.



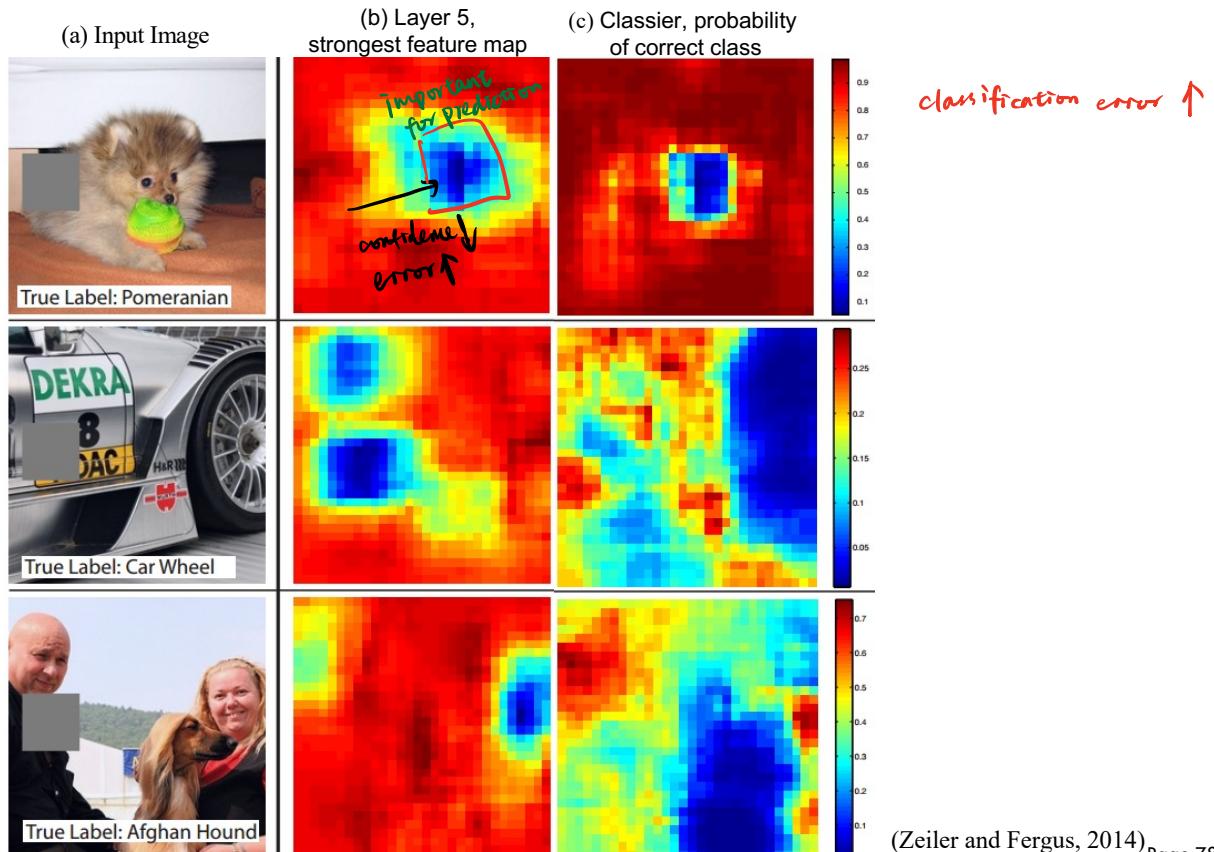
Visualize features

- Layer 2 responds to corners and other edge/color conjunctions.
- Layer 3 has more complex invariances, capturing similar textures (e.g. mesh patterns).
- Layer 4 shows significant variation, but is more class-specific.
- Layer 5 shows entire objects with significant pose variation.



Occlusion Experiments

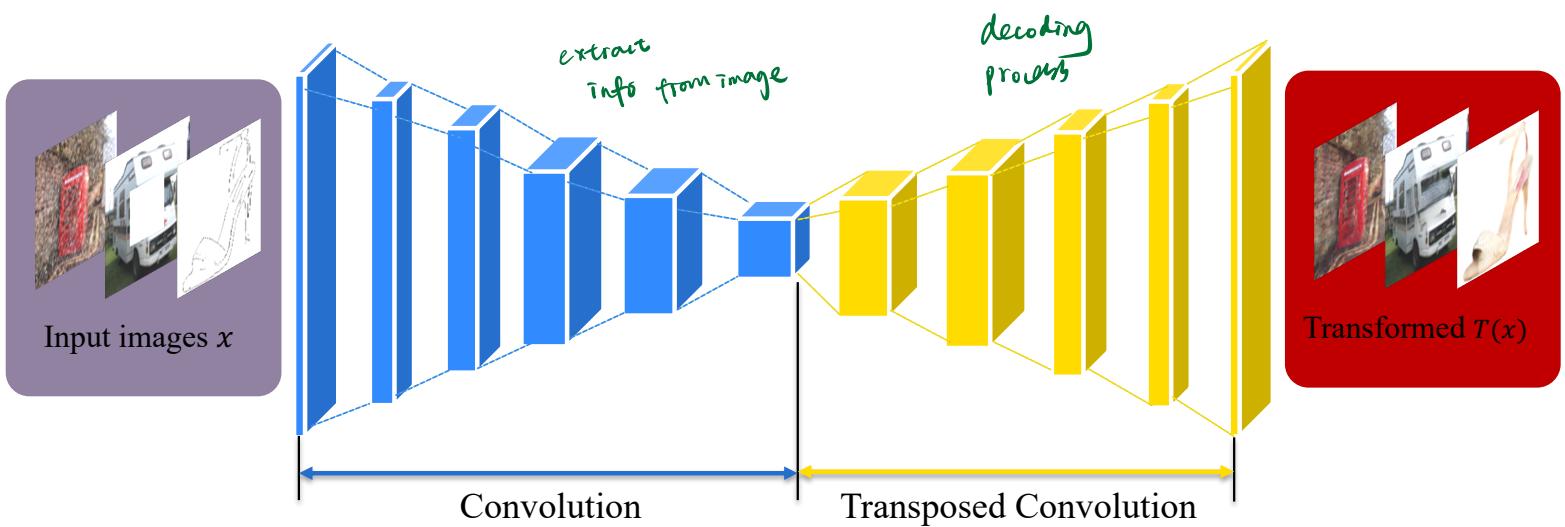
- Occlude images at different locations and visualize feature activations and classifier confidence



Variations of CNNs

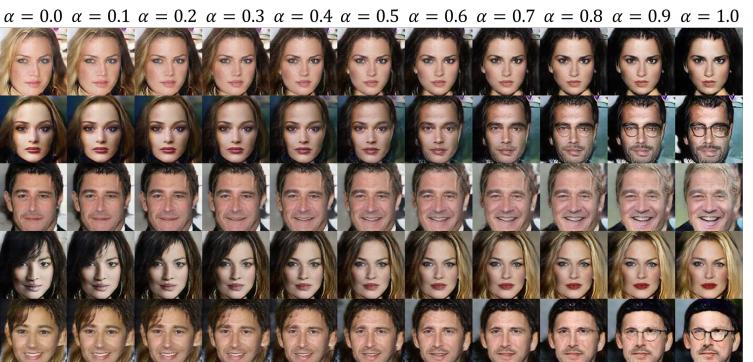
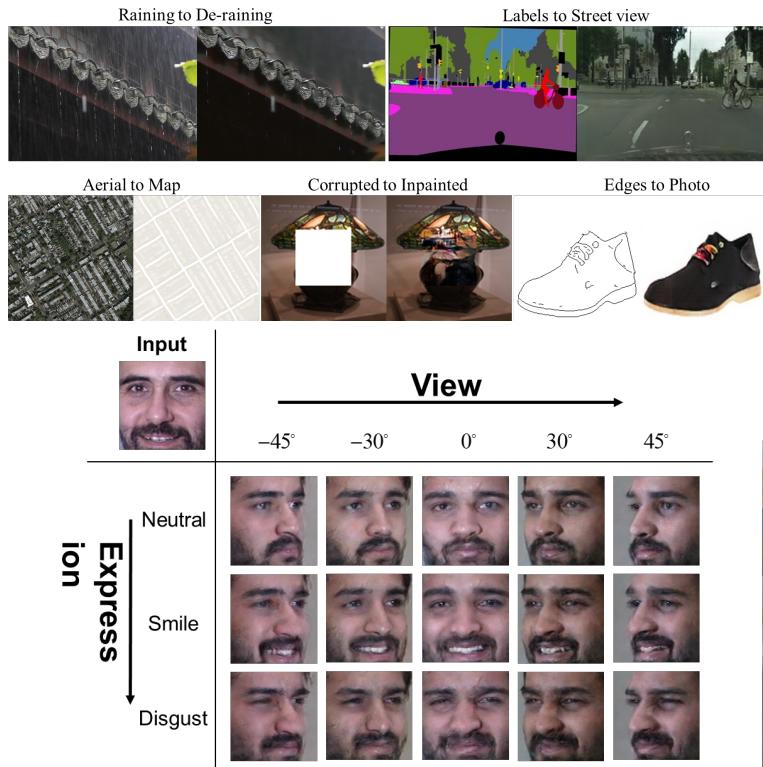
Transposed Convolution

- Transposed Convolution
 - Also named “Deconvolution”, “up-convolution” ...
 - Usually, the output is an image
 - Different from the above ‘Deconvolutional method’ for visualizing CNNs learned features



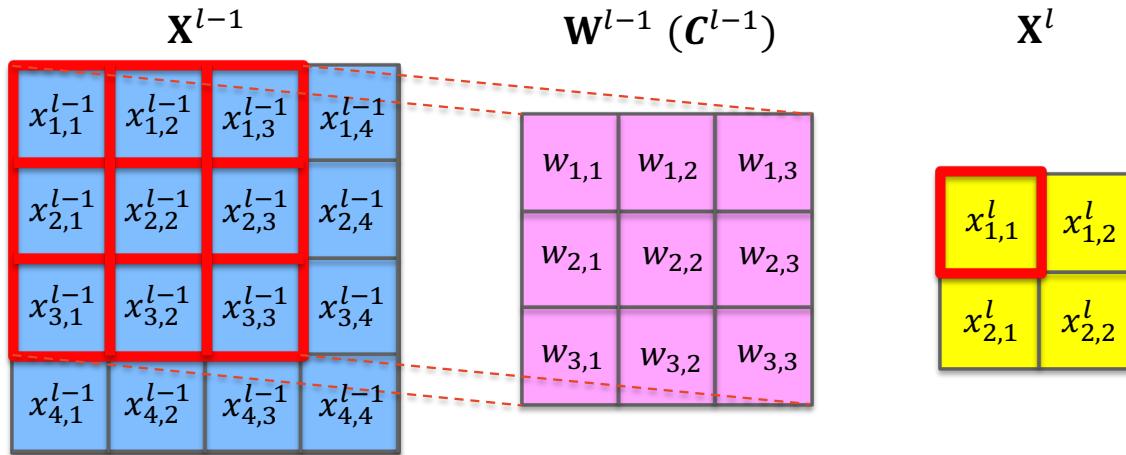
Transposed Convolution

- Applications
 - Generative model, encoder-decoder, image editing, image transformation ...



Transposed Convolution

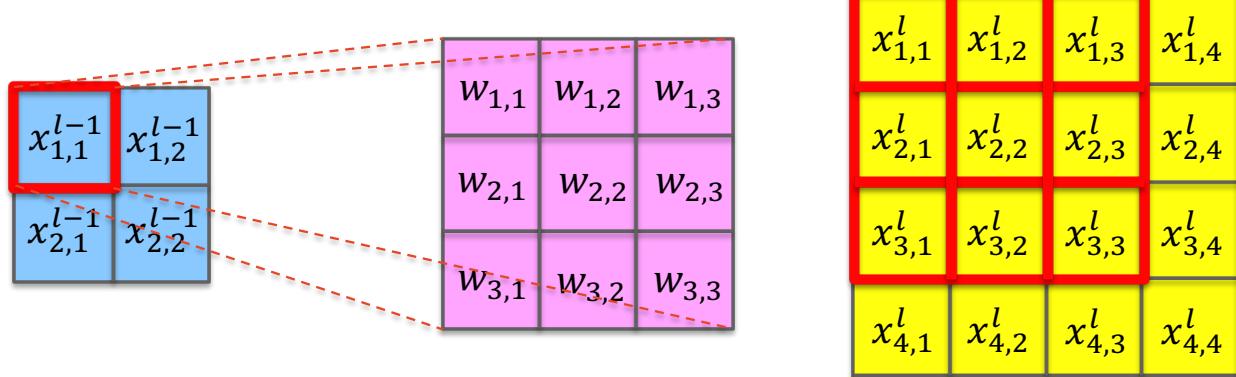
- Compare to aforementioned convolution operation, can we go the other way around, i.e., map from a 4-dim space to a 16-dim space, while keeping the connectivity pattern of the original convolution.



Convolution operation

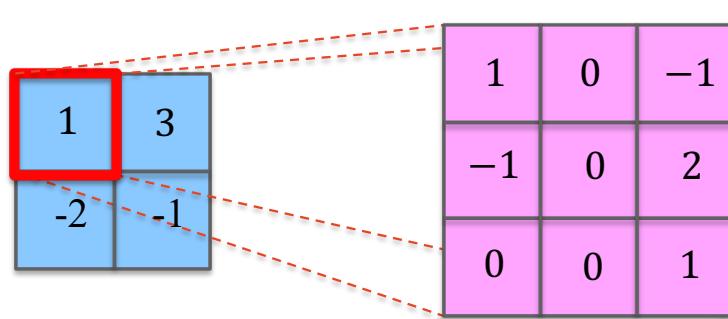
Transposed Convolution

- Compare to aforementioned convolution operation, can we go the other way around, i.e., map from a 4-dim space to a 16-dim space, while keeping the connectivity pattern of the original convolution.



Transposed Convolution

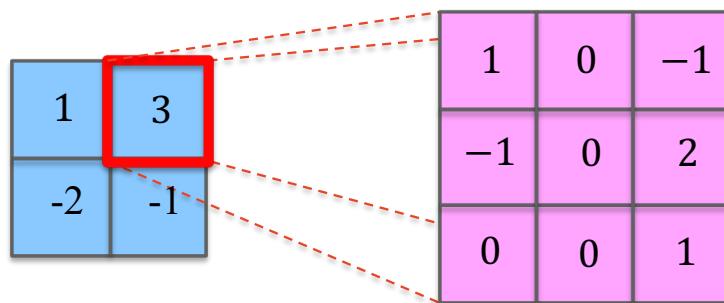
Transposed Convolution



1	0	-1	
-1	0	2	
0	0	1	

1	0	-1	
-1	0	2	
0	0	1	

Transposed Convolution

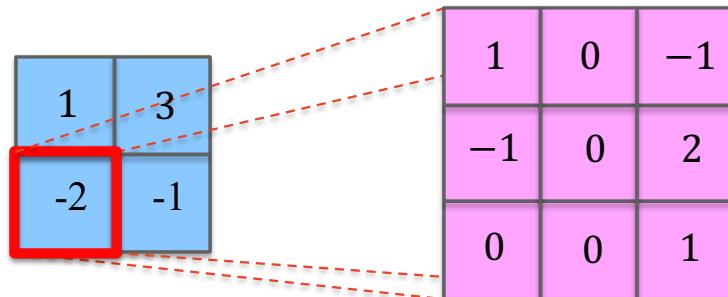


1	0	-1	-3
-1	-3	2	6
0	0	1	3

1	0	-1	
-1	0	2	
0	0	1	

	3	0	-3
	-3	0	6
	0	0	3

Transposed Convolution



1	3	-1	-3
-3	-3	4	6
0	0	-13	3
0	0	-1	

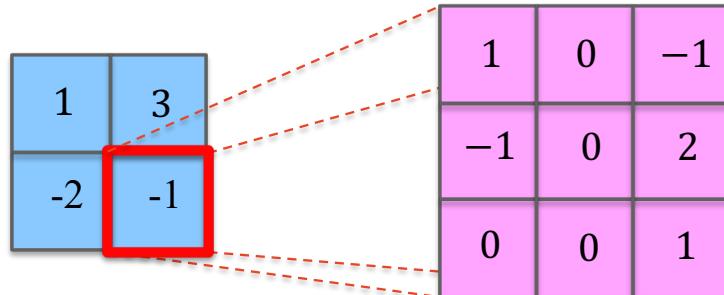
1	0	-1	
-1	0	2	
0	0	1	

	3	0	-3
	-3	0	6
	0	0	3

-2	0	2	
2	0	-4	
0	0	-2	

Transposed Convolution

overlap → use Σ



1	3	-1	-3
-3	-4	4	6
2	0	-3	3
0	0	-1	2

1	0	-1	
-1	0	2	
0	0	1	

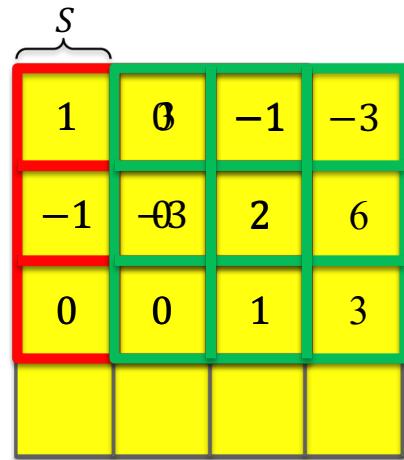
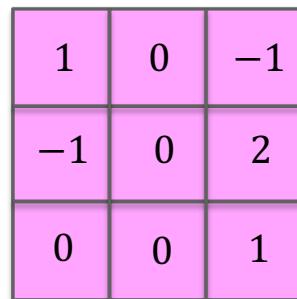
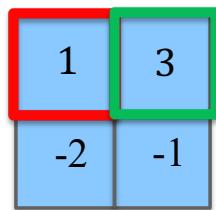
	3	0	-3
	-3	0	6
	0	0	3

-2	0	2	
2	0	-4	
0	0	-2	

-1	0	1	
1	0	-2	
0	0	-1	

Transposed Convolution

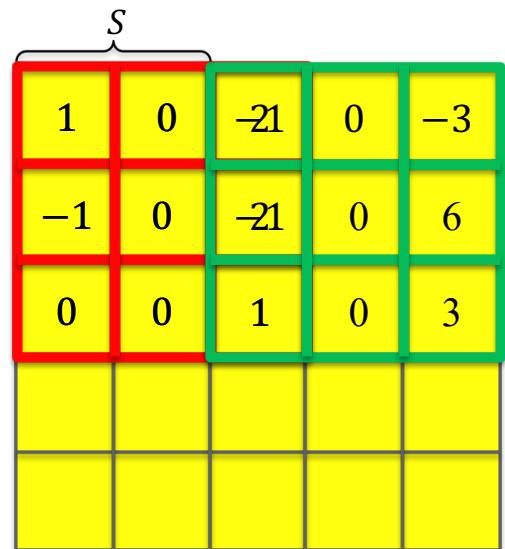
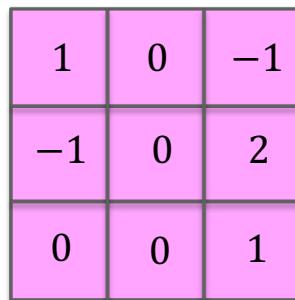
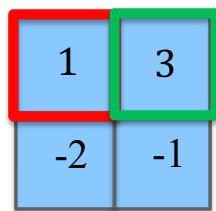
- Stride



Stride = 1

Transposed Convolution

- Stride



Stride = 2

Transposed Convolution

- Cropping

1	3
-2	-1

1	0	-1
-1	0	2
0	0	1

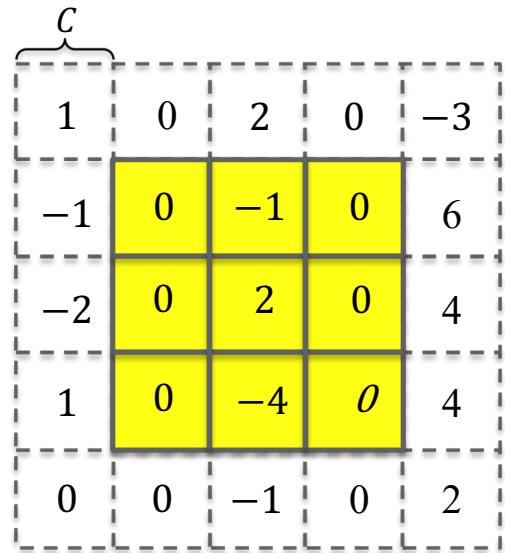
1	0	2	0	-3
-1	0	-1	0	6
-2	0	2	0	4
1	0	-4	0	4
0	0	-1	0	2

Transposed Convolution

- Cropping

1	3
-2	-1

1	0	-1
-1	0	2
0	0	1

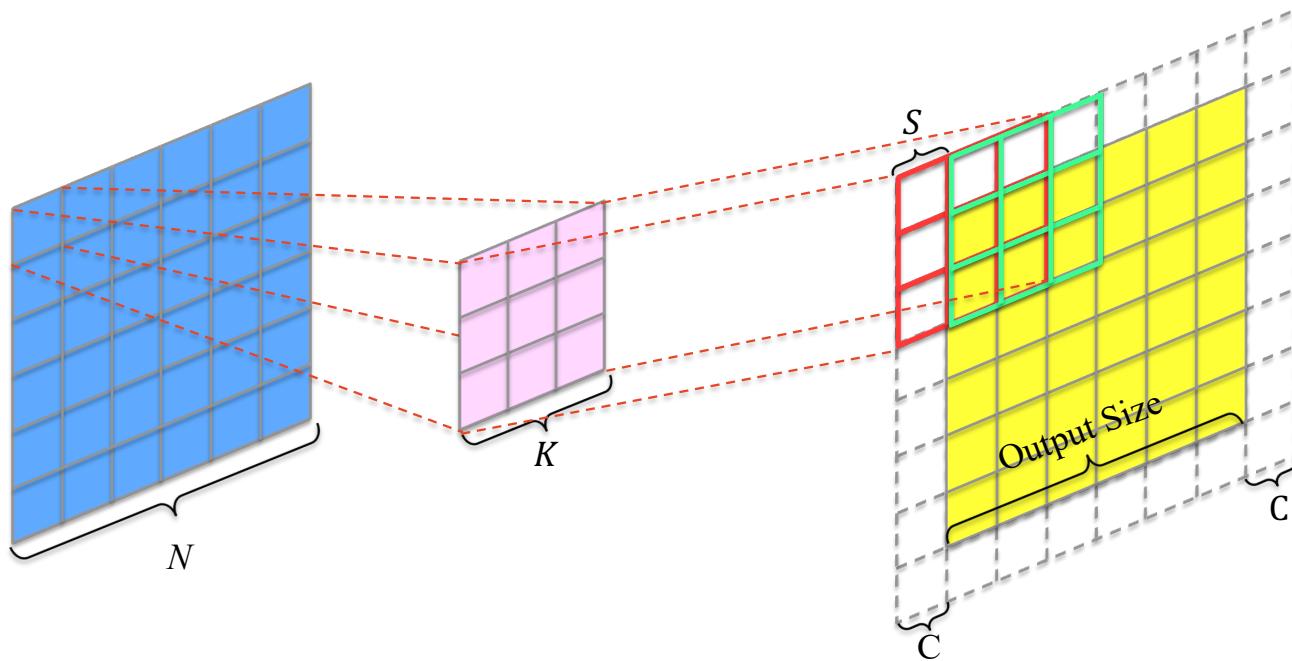


$$\text{Crop} = 1 \quad \leftrightarrow \quad \text{padding}$$

Transposed Convolution

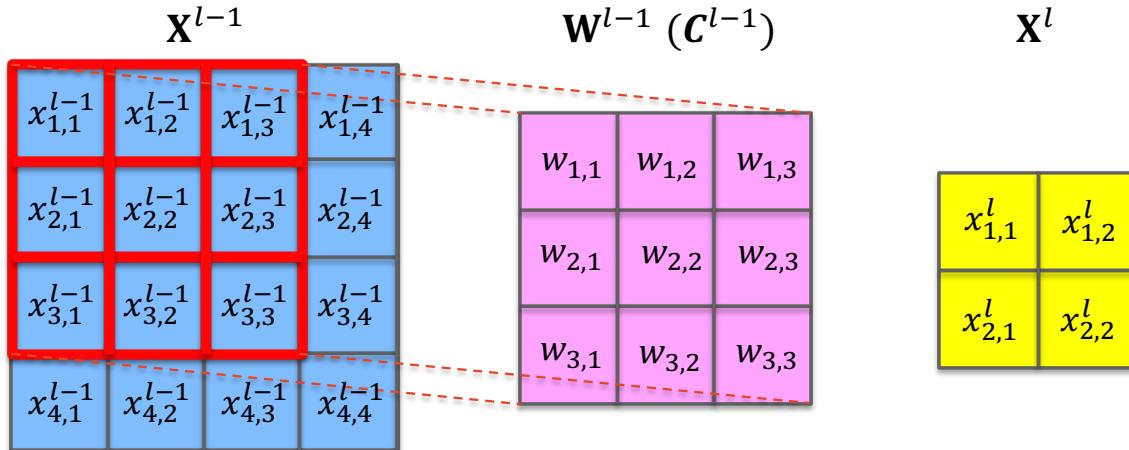
- Output Size

$$\text{Output Size} = (N - 1)S + K - 2C$$



Convolution as a matrix operation

- Recall the matrix operation of original convolution



$$\mathcal{C}^{l-1} \times (x_{1,1}^{l-1}, \dots, x_{4,4}^{l-1})^\top = (x_{1,1}^l, \dots, x_{2,2}^l)^\top$$

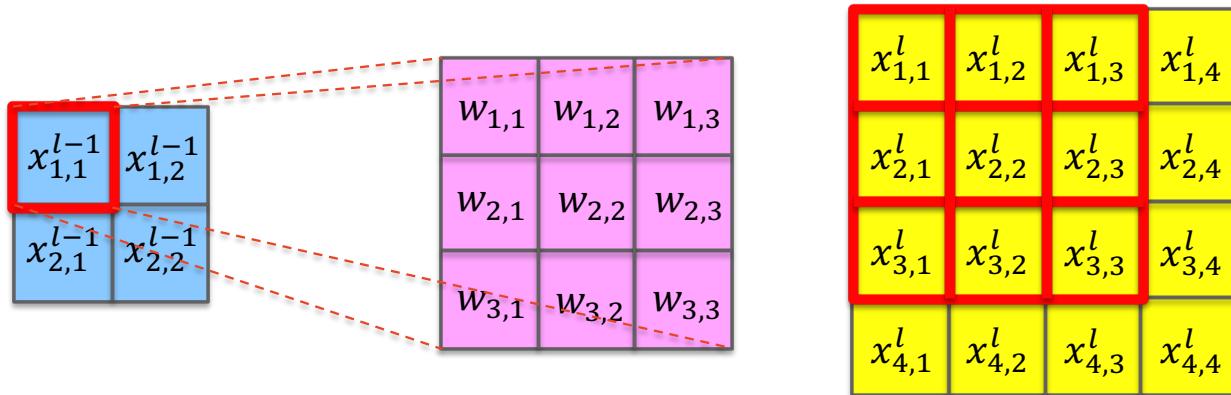
$$\mathcal{C}^{l-1} = \begin{pmatrix} w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} & 0 & 0 & 0 & 0 & 0 \\ 0 & w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} & 0 \\ 0 & 0 & 0 & 0 & 0 & w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} \end{pmatrix}$$

Transposed Convolution as a matrix operation

- Given the same matrix C , transposed convolution can be represented as:

$$C^\top \times (x_{1,1}^{l-1}, \dots, x_{2,2}^{l-1})^\top = (x_{1,1}^l, \dots, x_{4,4}^l)^\top$$

Forward pass of transposed convolution

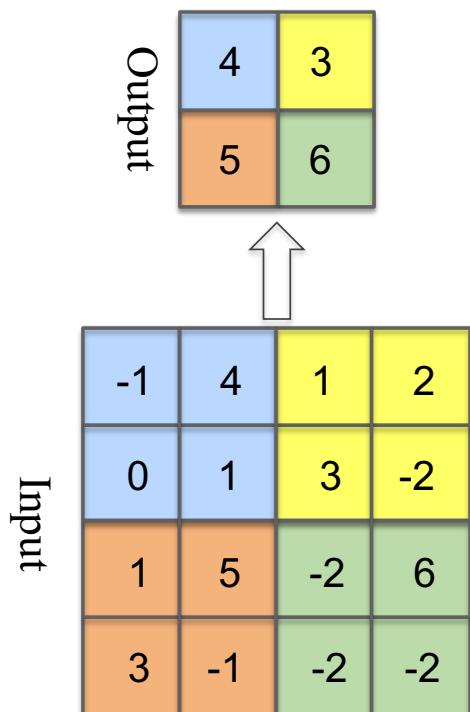


$$C = \begin{pmatrix} w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} & 0 & 0 & 0 & 0 & 0 \\ 0 & w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} & 0 \\ 0 & 0 & 0 & 0 & 0 & w_{1,1} & w_{1,2} & w_{1,3} & 0 & w_{2,1} & w_{2,2} & w_{2,3} & 0 & w_{3,1} & w_{3,2} & w_{3,3} \end{pmatrix}$$

Max pooling & Up-pooling

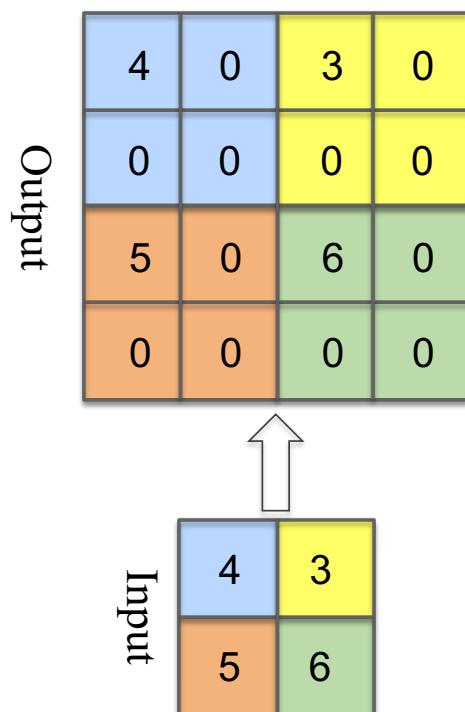
Max pooling

- Pooling ops: $\max(\cdot)$



Up-pooling

- Fill numbers to pre-defined position, others fill zeros.



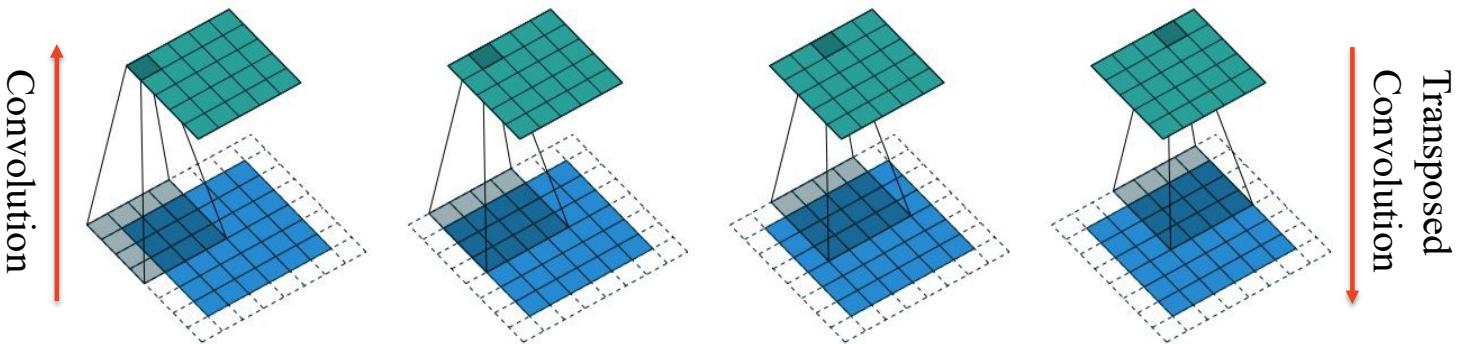
Transposed Convolution

Convolution

- From image to features
- Kernel: $n \times m \rightarrow 1$
- Padding
- ... → Filtering → Activate function → Pooling → ...

Transposed convolution

- From features to image
- Kernel: $1 \rightarrow n \times m$
- Cropping
- ... → Filtering → Activate function → Up-pooling → ...



https://github.com/vdumoulin/conv_arithmetic

Thank you!